

ADS-503 Team 4 Final Project

Leonard Littleton

Lina Nguyen

Emanuel Lucban

6/13/21

```
library(readr)
library(data.table)
library(mlbench)
library(ggplot2)
library(tidyr)
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
library(e1071)
library(caret)
```

```
## Loading required package: lattice
```

```
library(naniar)
library(MLmetrics)
```

```
##
```

```
## Attaching package: 'MLmetrics'
```

```
## The following objects are masked from 'package:caret':
```

```
##
```

```
##      MAE, RMSE
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
##      Recall
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:data.table':
```

```
##
```

```
##      between, first, last
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library(MASS)
```

```
##
```

```
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##      select

library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##      cov, smooth, var

# Parallel Processing
library(doParallel)

## Loading required package: foreach
## Loading required package: iterators
## Loading required package: parallel
cl <- makePSOCKcluster(8)
```

Synthetic Financial Datasets For Fraud Detection

Data Preparation

```
synth_df <- fread('./data/PS_20174392719_1491204439457_log.csv', header = TRUE)
synth_df <- subset(synth_df, select = -c(isFlaggedFraud, nameOrig, nameDest))

synth_df <- cbind(synth_df, data.frame(hours_intraday = synth_df$step % 24,
                                     by_day = round(synth_df$step / 24),
                                     day_of_week = round(synth_df$step / 24) % 7))

# Log transform (amount, oldbalanceOrg, newbalanceOrig, oldbalanceDest, newbalanceDest)
cont_vars <- c('amount', 'oldbalanceOrg', 'newbalanceOrig', 'oldbalanceDest', 'newbalanceDest')

# add small constant to prevent inf values
log_scaled <- sapply(data.frame(synth_df)[, cont_vars], function(x) log(x + 1))
colnames(log_scaled) <- lapply(cont_vars, function(x) paste('log_', x, sep=''))

synth_df <- cbind(synth_df, log_scaled)
synth_df$type <- as.factor(synth_df$type)
dmy <- dummyVars(" ~ type", data = synth_df, sep = '.', fullRank = TRUE)
synth_df <- cbind(synth_df, data.frame(predict(dmy, newdata = synth_df)))
# drop type
synth_df <- subset(synth_df, select = -c(type))
```

Split Data into Training and Test Datasets using Stratified Random Sampling

```
set.seed(42)
```

```

# split x and y
x <- subset(synth_df, select = -c(isFraud))
y <- synth_df$isFraud

data_part <- createDataPartition(y = y, p = 0.75, list = FALSE)

x_train <- x[data_part, ]
y_train <- y[data_part]
x_test <- x[-data_part, ]
y_test <- y[-data_part]

y_train <- as.factor(y_train)
y_test <- as.factor(y_test)
levels(y_train) <- c('no', 'yes')
levels(y_test) <- c('no', 'yes')

```

Modeling

```

# SET TRUE TO TRAIN MODELS
TRAIN = FALSE

# train control
ctrl <- trainControl(method = "cv",
                      number = 5,
                      summaryFunction = prSummary,
                      classProbs = TRUE,
                      savePredictions = TRUE)

# Register Parallel Processing
registerDoParallel(cl)

```

Neural Network

Training

```

if (TRAIN){
  nnTuning <- expand.grid(size = c(1:5),
                        decay = c(0.01, 0.05, 0.1))
  nnetModel <- train(x_train, y_train,
                    method = "nnet",
                    metric = 'AUC',
                    trControl= ctrl,
                    tuneGrid = nnTuning,
                    preProcess=c("scale","center"))

  saveRDS(nnetModel, "rds_files/nnet_model.rds")
} else {
  nnetModel <- readRDS('rds_files/nnet_model.rds')
}

```

```
nnetModel
```

```
## Neural Network
##
```

```
## 4771965 samples
##      18 predictor
##      2 classes: 'no', 'yes'
##
## Pre-processing: scaled (18), centered (18)
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 3817572, 3817571, 3817572, 3817572, 3817573
## Resampling results across tuning parameters:
##
##  size  decay  AUC          Precision  Recall    F
##  1     0.01  0.03841136  0.9991485  0.9998995  0.9995237
##  1     0.05  0.08480826  0.9995132  0.9999060  0.9997095
##  1     0.10  0.05056086  0.9991267  0.9999748  0.9995505
##  2     0.01  0.00000000  0.9987188  1.0000000  0.9993590
##  2     0.05  0.46270579  0.9993290  0.9999692  0.9996489
##  2     0.10  0.20514610  0.9991269  0.9999287  0.9995275
##  3     0.01  0.15922016  0.9987188  1.0000000  0.9993590
##  3     0.05  0.11514449  0.9989077  0.9999973  0.9994522
##  3     0.10  0.07911778  0.9989279  0.9999914  0.9994593
##  4     0.01  0.00000000  0.9987188  1.0000000  0.9993590
##  4     0.05  0.03219574  0.9987188  1.0000000  0.9993590
##  4     0.10  0.00000000  0.9987188  1.0000000  0.9993590
##  5     0.01  0.00000000  0.9987188  1.0000000  0.9993590
##  5     0.05  0.10122763  0.9988911  0.9999713  0.9994309
##  5     0.10  0.00000000  0.9987188  1.0000000  0.9993590
##
## AUC was used to select the optimal model using the largest value.
## The final values used for the model were size = 2 and decay = 0.05.
```

Neural Network Prediction on Test Data

```
nnetPred <- predict(nnetModel, newdata = x_test)
nnetCFM <- confusionMatrix(nnetPred, y_test, positive = 'yes')
nnetCFM
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      no      yes
##           no 1588556    2099
##           yes      0         0
##
##           Accuracy : 0.9987
##           95% CI : (0.9986, 0.9987)
##           No Information Rate : 0.9987
##           P-Value [Acc > NIR] : 0.5058
##
##           Kappa : 0
##
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.00000
##           Specificity : 1.00000
##           Pos Pred Value :      NaN
##           Neg Pred Value : 0.99868
```

```
##           Prevalence : 0.00132
##           Detection Rate : 0.00000
##           Detection Prevalence : 0.00000
##           Balanced Accuracy : 0.50000
##
##           'Positive' Class : yes
##
```

Neural Network Variable Importance

```
varImp(nnetModel)
```

```
## nnet variable importance
##
##           Overall
## newbalanceOrig    100.000
## oldbalanceOrig     94.576
## type.DEBIT         55.842
## oldbalanceDest     45.186
## log_newbalanceOrig 43.786
## log_oldbalanceDest 41.047
## type.CASH_OUT      39.638
## log_oldbalanceOrg  32.570
## log_newbalanceDest 31.992
## type.TRANSFER       26.864
## type.PAYMENT        25.894
## newbalanceDest      23.269
## log_amount          20.260
## amount              19.157
## day_of_week         18.376
## by_day              13.672
## step                5.901
## hours_intraday       0.000
```

Linear Discriminant Analysis

Training

```
if (TRAIN) {
  ldaModel <- train(x_train, y_train,
                    method = "lda",
                    metric = 'AUC',
                    trControl= ctrl,
                    preProcess=c("scale","center"),
                    tuneLength = 10,
                    verbose = FALSE, trace = FALSE)

  saveRDS(ldaModel, "rds_files/lda_model.rds")
} else {
  ldaModel <- readRDS('rds_files/lda_model.rds')
}
```

```
ldaModel
```

```
## Linear Discriminant Analysis
```

```
##
## 4771965 samples
##      18 predictor
##      2 classes: 'no', 'yes'
##
## Pre-processing: scaled (18), centered (18)
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 3817572, 3817571, 3817572, 3817572, 3817573
## Resampling results:
##
##      AUC          Precision  Recall    F
##  0.0005841567  0.9992177  0.9999041  0.9995608
```

LDA Prediction on Test Data

```
ldaPred <- predict(ldaModel, newdata = x_test)
ldaCFM <- confusionMatrix(ldaPred, y_test, positive = 'yes')
ldaCFM
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction      no      yes
##           no 1588403    1253
##           yes   153     846
##
##              Accuracy : 0.9991
##              95% CI : (0.9991, 0.9992)
##      No Information Rate : 0.9987
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.5458
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.4030491
##              Specificity : 0.9999037
##              Pos Pred Value : 0.8468468
##              Neg Pred Value : 0.9992118
##              Prevalence : 0.0013196
##              Detection Rate : 0.0005319
##      Detection Prevalence : 0.0006280
##              Balanced Accuracy : 0.7014764
##
##              'Positive' Class : yes
##
```

LDA Variable Importance

```
varImp(ldaModel)

## ROC curve variable importance
##
##              Importance
## log_oldbalanceOrg    100.00
## oldbalanceOrg        100.00
## log_amount           93.07
```

```
## amount                93.07
## type.TRANSFER          66.00
## newbalanceOrig        65.53
## log_newbalanceOrig     65.53
## type.PAYMENT           54.10
## step                   54.00
## by_day                 52.76
## hours_intraday         49.38
## log_oldbalanceDest     41.67
## oldbalanceDest        41.67
## type.CASH_OUT          23.98
## day_of_week            15.74
## newbalanceDest         11.49
## log_newbalanceDest     11.49
## type.DEBIT             0.00
```

Quadratic Discriminant Analysis

Training

```
if (TRAIN) {
  qdaModel <- train(x_train, y_train,
                    method = "qda",
                    metric = 'AUC',
                    trControl= ctrl,
                    preProcess=c("pca", "scale","center"),
                    tuneLength = 10,
                    verbose = FALSE, trace = FALSE)
  saveRDS(qdaModel, "rds_files/qda_model.rds")
} else {
  qdaModel <- readRDS('rds_files/qda_model.rds')
}
```

```
qdaModel
```

```
## Quadratic Discriminant Analysis
##
## 4771965 samples
##      18 predictor
##      2 classes: 'no', 'yes'
##
## Pre-processing: principal component signal extraction (18), scaled
## (18), centered (18)
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 3817572, 3817571, 3817572, 3817572, 3817573
## Resampling results:
##
##      AUC          Precision  Recall      F
##  0.2011124  0.9997956  0.9934996  0.9966375
```

QDA Prediction on Test Data

```
qdaPred <- predict(qdaModel, newdata = x_test)
qdaCFM <- confusionMatrix(qdaPred, y_test, positive = 'yes')
qdaCFM
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      no      yes
##           no 1578607      295
##           yes   9949     1804
##
##           Accuracy : 0.9936
##           95% CI : (0.9934, 0.9937)
##           No Information Rate : 0.9987
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.2588
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.859457
##           Specificity : 0.993737
##           Pos Pred Value : 0.153493
##           Neg Pred Value : 0.999813
##           Prevalence : 0.001320
##           Detection Rate : 0.001134
##           Detection Prevalence : 0.007389
##           Balanced Accuracy : 0.926597
##
##           'Positive' Class : yes
##
```

QDA Variable Importance

```
varImp(qdaModel)
```

```
## ROC curve variable importance
##
##           Importance
## log_oldbalanceOrg      100.00
## oldbalanceOrg          100.00
## log_amount             93.07
## amount                 93.07
## type.TRANSFER          66.00
## newbalanceOrig         65.53
## log_newbalanceOrig     65.53
## type.PAYMENT           54.10
## step                  54.00
## by_day                 52.76
## hours_intraday         49.38
## log_oldbalanceDest     41.67
## oldbalanceDest        41.67
## type.CASH_OUT          23.98
## day_of_week            15.74
## newbalanceDest         11.49
## log_newbalanceDest     11.49
## type.DEBIT             0.00
```



```
# Stop cluster and parallel processing  
stopCluster(cl)  
registerDoSEQ()
```